

## Import Libraries

In [ ]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
import plotly.graph_objects as go
from scipy import stats
import os,sys
from scipy.stats import norm, skew
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.linear_model import LinearRegression
import warnings
warnings.filterwarnings('ignore')
```

## Import data

In [ ]:

```
df_train=pd.read_csv('Train.csv')
df_test=pd.read_csv('Test.csv')
```

## Data Understanding

In [ ]:

```
df_train.head()
```

Out[3]:

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier
0	FDA15	9.30	Low Fat	0.016047	Dairy	249.8092	OUT0.
1	DRC01	5.92	Regular	0.019278	Soft Drinks	48.2692	OUT0
2	FDN15	17.50	Low Fat	0.016760	Meat	141.6180	OUT0.
3	FDX07	19.20	Regular	0.000000	Fruits and Vegetables	182.0950	OUT0
4	NCD19	8.93	Low Fat	0.000000	Household	53.8614	OUT0

In [ ]:

```
df_train.dtypes
```

Out[4]:

```
Item_Identifier      object
Item_Weight          float64
Item_Fat_Content      object
Item_Visibility      float64
Item_Type            object
Item_MRP             float64
Outlet_Identifier     object
Outlet_Establishment_Year  int64
Outlet_Size          object
Outlet_Location_Type  object
Outlet_Type          object
Item_Outlet_Sales     float64
dtype: object
```

In [ ]:

```
df_train.shape
```

Out[5]:

```
(8523, 12)
```

In [ ]:

```
df_train.isna().any()
```


Out[6]:

```
Item_Identifier      False
Item_Weight          True
Item_Fat_Content      False
Item_Visibility      False
Item_Type            False
Item_MRP             False
Outlet_Identifier     False
Outlet_Establishment_Year  False
Outlet_Size          True
Outlet_Location_Type  False
Outlet_Type          False
Item_Outlet_Sales     False
dtype: bool
```

## Data Visualization

In [ ]:

```
plt.figure(figsize=(10,8))
sns.heatmap(df_train.isna(),cmap='gnuplot')
plt.show()
```



yellow mark show null values

## 1. Item Weight

In [ ]:

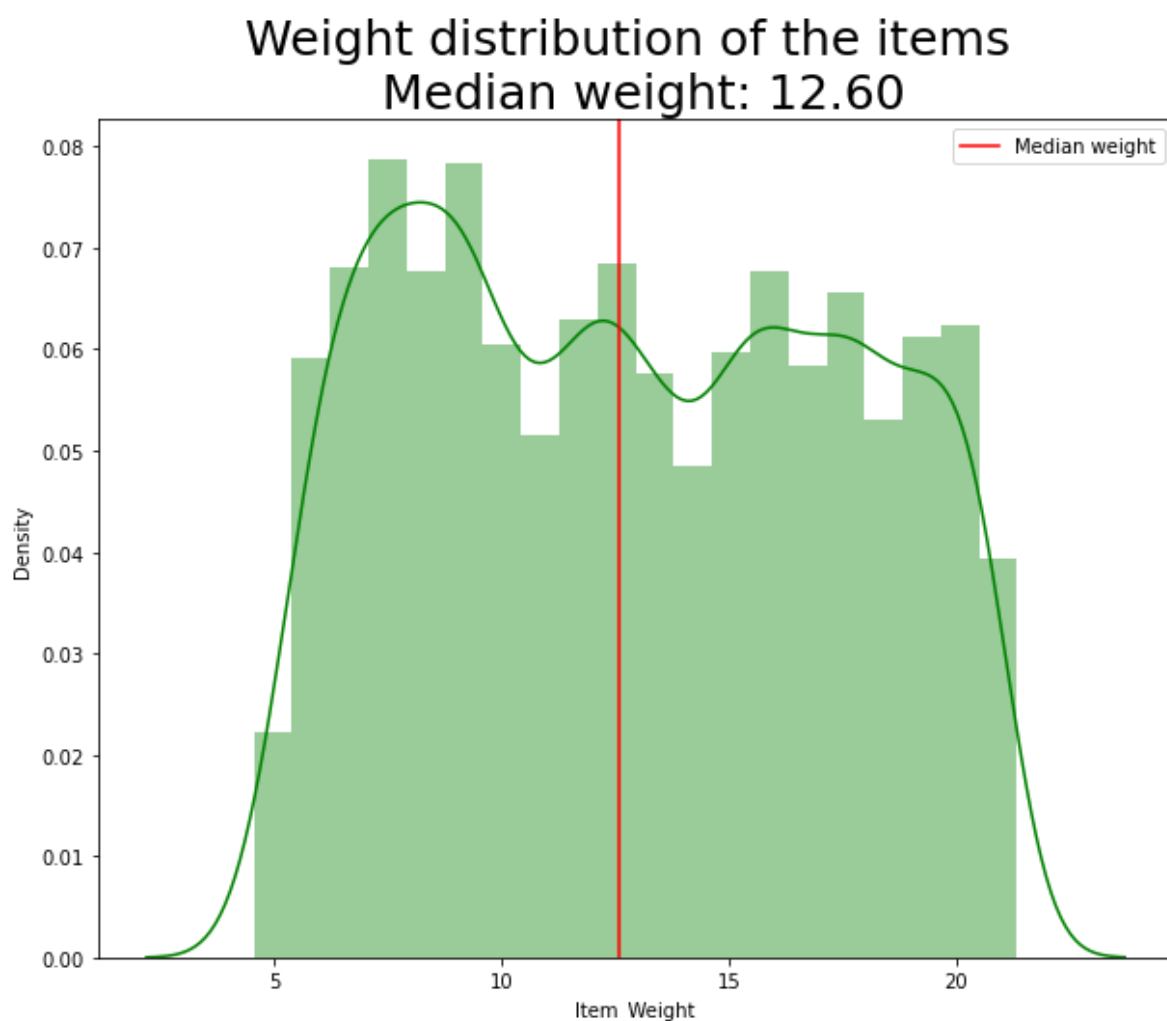
```
df_train['Item_Weight'].describe()
```

Out[8]:

```
count      7060.000000
mean       12.857645
std        4.643456
min        4.555000
25%        8.773750
50%       12.600000
75%       16.850000
max       21.350000
Name: Item_Weight, dtype: float64
```

In [ ]:

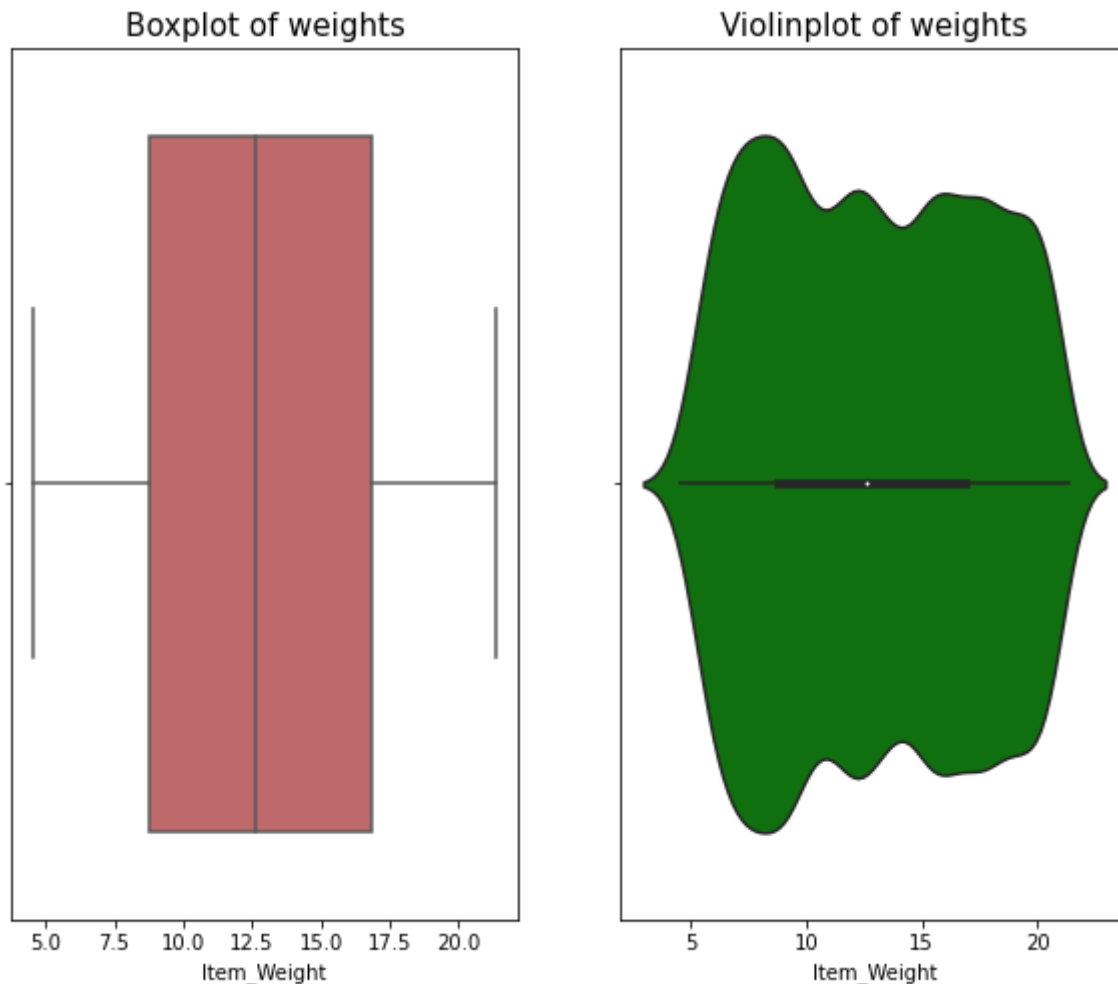
```
plt.figure(figsize=(10,8))
sns.distplot(df_train['Item_Weight'].dropna(),color='green')
plt.title('Weight distribution of the items \n Median weight: {0:.2f}'.format(df_train['Item_Weight'].dropna().median()))
plt.axvline(df_train['Item_Weight'].dropna().median(),color='red',label='Median weight')
plt.legend()
plt.show()
```



In [ ]:

```
fig1=plt.figure(figsize=(10,8))
ax1=fig1.add_subplot(121)
sns.boxplot(df_train['Item_Weight'],ax=ax1,orient='v',color='indianred')
ax1.set_title('Boxplot of weights',size=15)

ax2=fig1.add_subplot(122)
sns.violinplot(df_train['Item_Weight'],ax=ax2,orient='v',color='green')
ax2.set_title('Violinplot of weights',size=15)
plt.show()
```



As we can see from the above violin and distplot, the curve plateaus over a large range of weights. Hence, it is simply not possible for us to assume a weight for the null values. We shall leave them as it is or drop them if it is later deemed to not be too important in our analysis.

## 2. Item Fat Content

In [ ]:

```
df_train['Item_Fat_Content'].unique()
```

Out[11]:

```
array(['Low Fat', 'Regular', 'low fat', 'LF', 'reg'], dtype=object)
```

In [ ]:

```
df_train['Item_Fat_Content']=df_train['Item_Fat_Content'].replace('low fat','Low Fat')  
df_train['Item_Fat_Content']=df_train['Item_Fat_Content'].replace('LF','Low Fat')  
df_train['Item_Fat_Content']=df_train['Item_Fat_Content'].replace('reg','Regular')  
df_train['Item_Fat_Content'].unique()
```

Out[12]:

```
array(['Low Fat', 'Regular'], dtype=object)
```

In [ ]:

```
df_train['Count']=1
df_fat=df_train.groupby('Item_Fat_Content')['Count'].sum().reset_index()

fig2=px.pie(df_fat,values='Count',names='Item_Fat_Content',hole=0.4)

fig2.update_layout(title='Fat content',title_x=0.48,
                    annotations=[dict(text='Fat',font_size=15, showarrow=False,height=800,width=100)],
fig2.update_traces(textfont_size=15,textinfo='percent+label')
fig2.show()
```

### 3. Item Visibility

In [ ]:

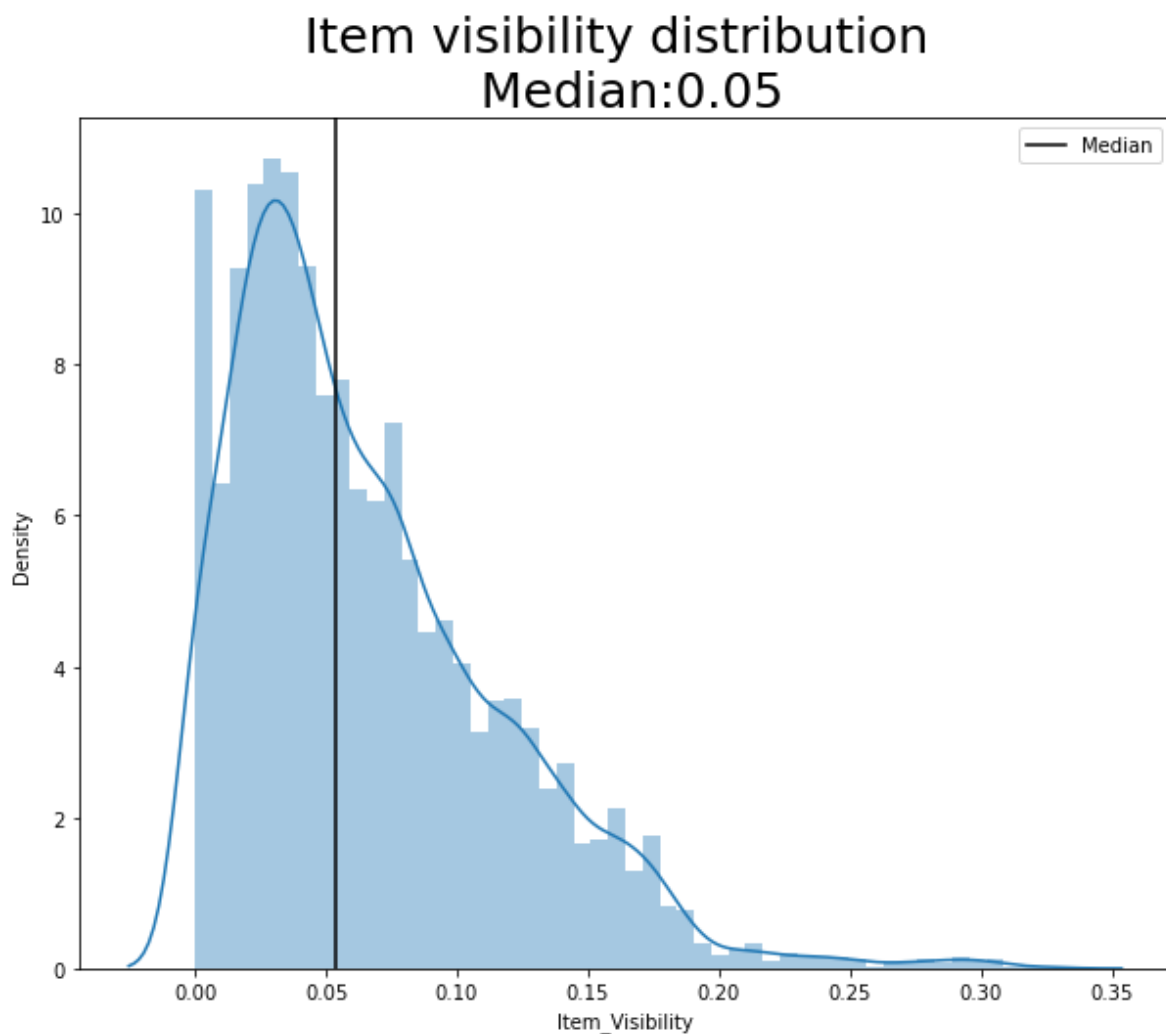
```
df_train['Item_Visibility'].describe()
```

Out[14]:

```
count      8523.000000
mean        0.066132
std         0.051598
min         0.000000
25%         0.026989
50%         0.053931
75%         0.094585
max         0.328391
Name: Item_Visibility, dtype: float64
```

In [ ]:

```
plt.figure(figsize=(10,8))
sns.distplot(df_train['Item_Visibility'])
plt.title('Item visibility distribution \n Median:{0:.2f}'.format(df_train['Item_Visibility'].median()))
plt.axvline(df_train['Item_Visibility'].median(),color='black',label='Median')
plt.legend()
plt.show()
```



we can see that it is right skew curve. Hence, a median would give us better indication than a mean value.




In [ ]:

```

sns.set()
fig3=plt.figure(figsize=(10,5))
ax1=fig3.add_subplot(121)
sns.boxplot(df_train['Item_Visibility'],orient='v',ax=ax1,color='green')
ax2=fig3.add_subplot(122)
stats.probplot(df_train['Item_Visibility'],plot=ax2)
plt.show()

```



As we can see, values above 0.2 visibility are outliers. Presence of outliers don't bode well with machine learning algos. Hence, we need to remove the outliers and try to form a normal distribution.

The probplot also seems to suggest that the values are deviating from the normal values after 0.2

In [ ]:


```
df_train[df_train['Item_Visibility']>0.2].shape[0]
```

Out[17]:

134

In [ ]:

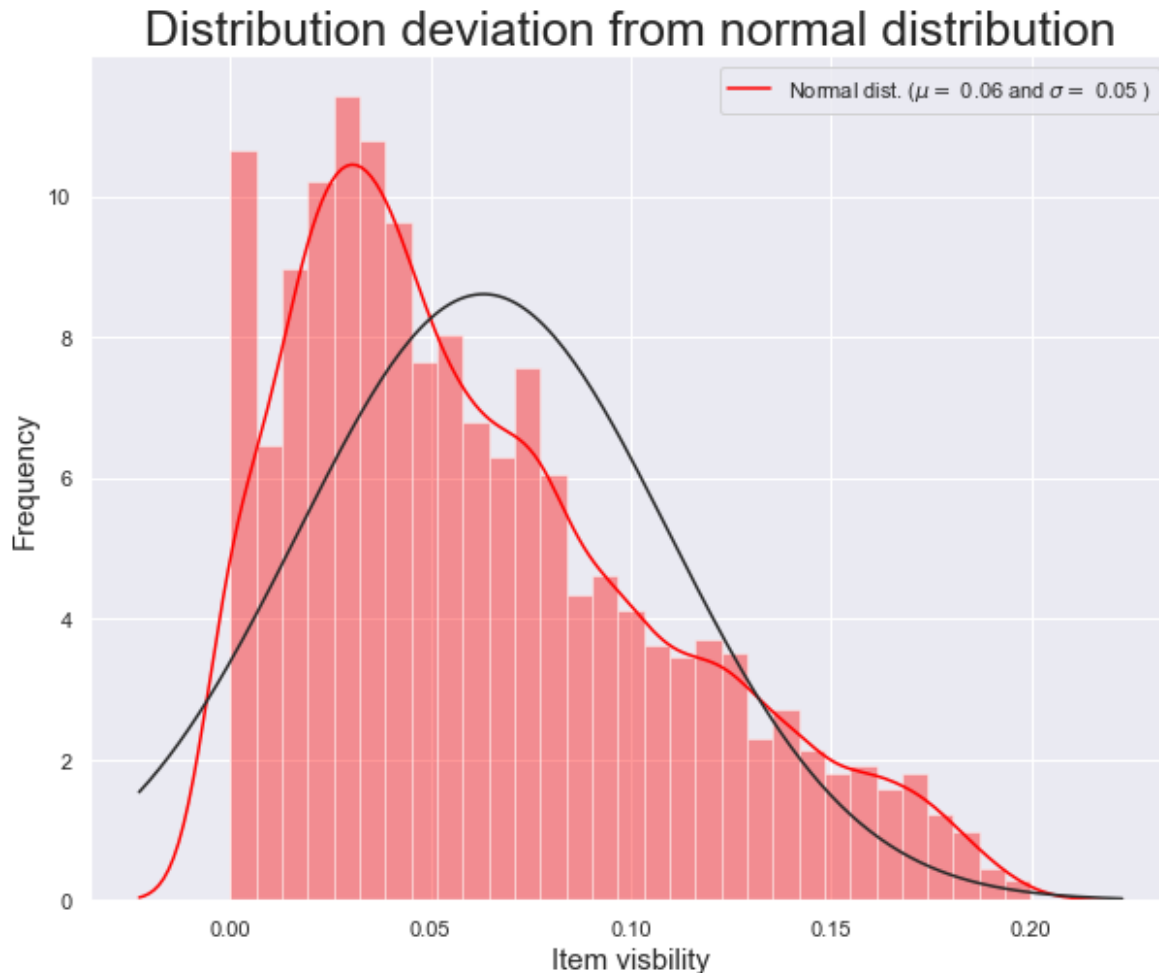
```
df_train=df_train[df_train['Item_Visibility']<0.2]  
stats.probplot(df_train['Item_Visibility'],plot=plt)  
plt.show()
```



Now, we see that the values above 0 are following a normal distribution to some extent.

In [ ]:

```
plt.figure(figsize=(10,8))
sns.distplot(df_train['Item_Visibility'],fit=norm,color='red')
plt.title('Distribution deviation from normal distribution',size=25)
plt.ylabel('Frequency',size=15)
plt.xlabel('Item visibility',size=15)
mu=df_train['Item_Visibility'].mean()
sigma=df_train['Item_Visibility'].std()
plt.legend(['Normal dist. ( $\mu = 0.06$  and  $\sigma = 0.05$ )'].format(mu, sigma)])
plt.show()
```



This is as close to a normal distribution I could get for our model

## 4. Item Type

In [ ]:


```
df_type=df_train.groupby('Item_Type')['Count'].sum().reset_index()
fig4=px.sunburst(df_train,path=['Item_Type','Item_Fat_Content'],names='Item_Type',color_con
fig4.update_layout(title='Item types',title_x=0.2,title_y=0.8,
                    annotations=[dict(showarrow=True,height=1000,width=900)],margin=dict(l=20
fig4.show()

fig5=px.pie(df_type,values='Count',names='Item_Type')
fig5.update_layout(title='Item distribution',title_x=0.1,title_y=0.8)
fig5.update_traces(textfont_size=15,textinfo='percent')
fig5.show()
```

## 5. Item Type

In [ ]:

```
plt.figure(figsize=(15,10))
sns.distplot(df_train['Item_MRP'],color='orange')
plt.title('Item MRP distribution \n Median:{0:.2f} Rs'.format(df_train['Item_MRP'].median()))
plt.axvline(df_train['Item_MRP'].median(),color='black',label='Median')
plt.legend()
plt.show()
```



As we can see, we don't have any clear distribution of the prices here. The distribution is multi modal in nature with multiple peaks.

The graph basically:


we have fair number of products whose prices range from 25-75 Rs. we have fair number of products in the 80-120 Rs range. Infact, it is the highest. the products increase again from 150-200 Rs range. There are fair number of products from 220-240 Rs range aswell.

In [ ]:

```
labels=df_train['Item_Type'].unique()
fig6=plt.figure(figsize=(10,10))
ax1=fig6.add_subplot(211)
sns.boxplot(x='Item_Type',y='Item_MRP',data=df_train,ax=ax1)
ax1.set_xticklabels(labels, rotation=75,size=9)

ax2=fig6.add_subplot(212)
sns.boxplot(x='Item_Fat_Content',y='Item_MRP',data=df_train,ax=ax2)

fig6.tight_layout(pad=3)
plt.show()
```



From the above plot, we see which item types have high MRPs. Dairy product and Starchy foods have a higher median price than the rest.

Both low and regular food have almost identical median price.

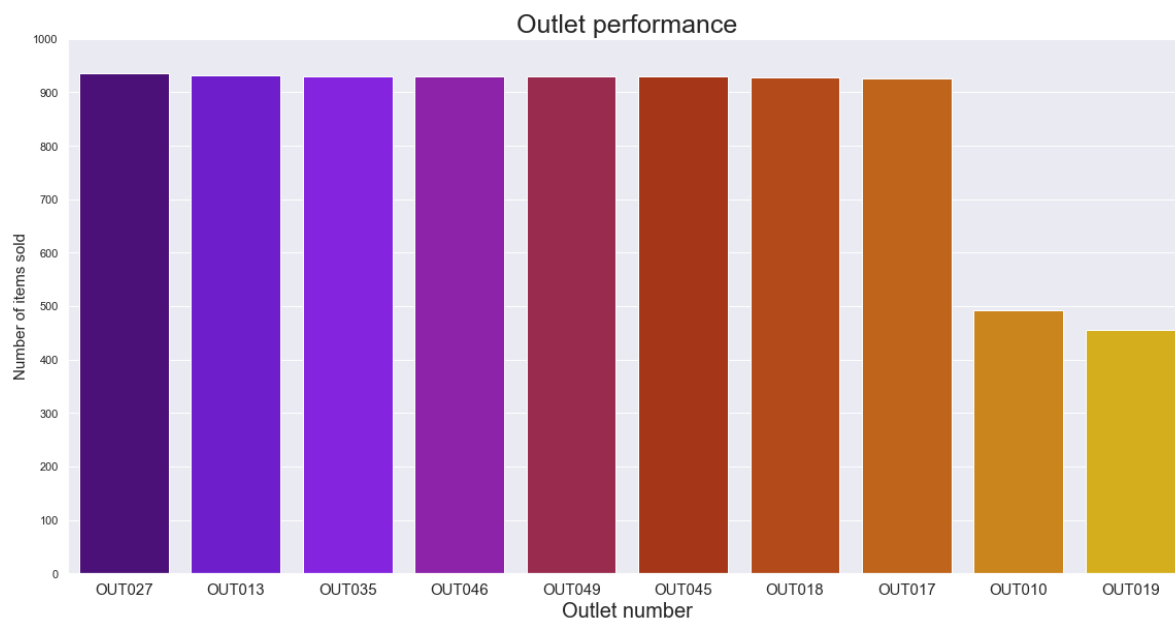
## 6. Outlet Identifier

In [ ]:

```
df_outlets=df_train.groupby('Outlet_Identifier')['Count'].sum().reset_index().sort_values(b
```

In [ ]:


```
sns.catplot('Outlet_Identifier', 'Count', data=df_outlets, aspect=2, height=8, kind='bar', palette='magma')
plt.xticks(size=15)
plt.ylabel('Number of items sold', size=15)
plt.xlabel('Outlet number', size=20)
plt.title('Outlet performance', size=25)
plt.yticks(np.arange(0, 1100, 100))
plt.show()
```





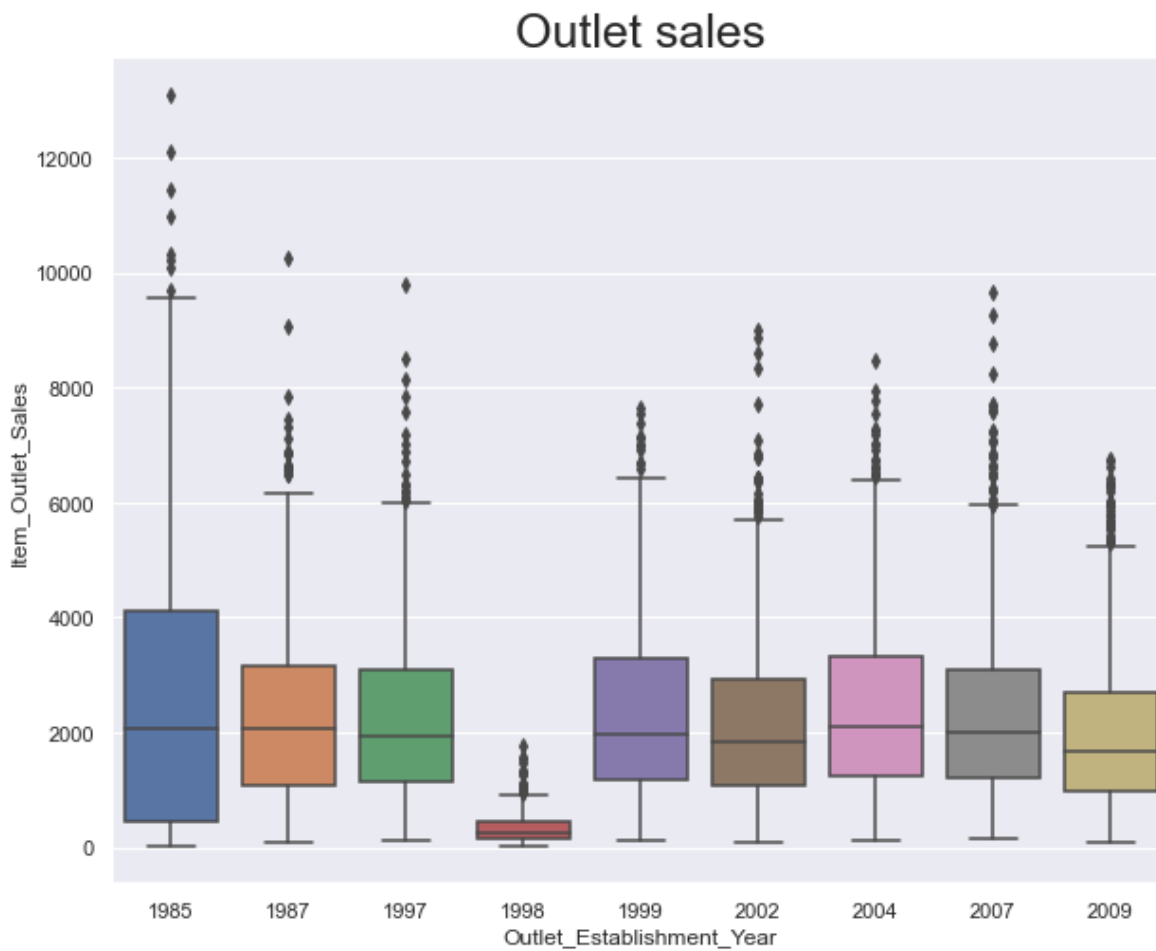
In [ ]:

```
plt.figure(figsize=(10,8))  
sns.boxplot('Outlet_Establishment_Year', 'Item_MRP', data=df_train)  
plt.show()
```



In [ ]:

```
plt.figure(figsize=(10,8))
sns.boxplot('Outlet_Establishment_Year', 'Item_Outlet_Sales', data=df_train)
plt.title('Outlet sales', size=25)
plt.show()
```



## 7. Outlet size

In [ ]:

```
df_train['Outlet_Size'].isna().value_counts()
```

Out[27]:

```
False    6041
True      2348
Name: Outlet_Size, dtype: int64
```

In [ ]:

```
df_size=df_train.groupby('Outlet_Size')['Count'].sum().reset_index()
fig7=px.pie(df_size,values='Count',names='Outlet_Size',hole=0.4)
fig7.update_layout(title='Store sizes',title_x=0.5,annotations=[dict(text='Fat',font_size=1
fig7.update_traces(textfont_size=15,textinfo='percent+label')
fig7.show()
```

In [ ]:

```
df_size_sales=df_train.groupby('Outlet_Size')[['Item_MRP','Item_Outlet_Sales']].mean().rese
df_size_sales
```

Out[29]:

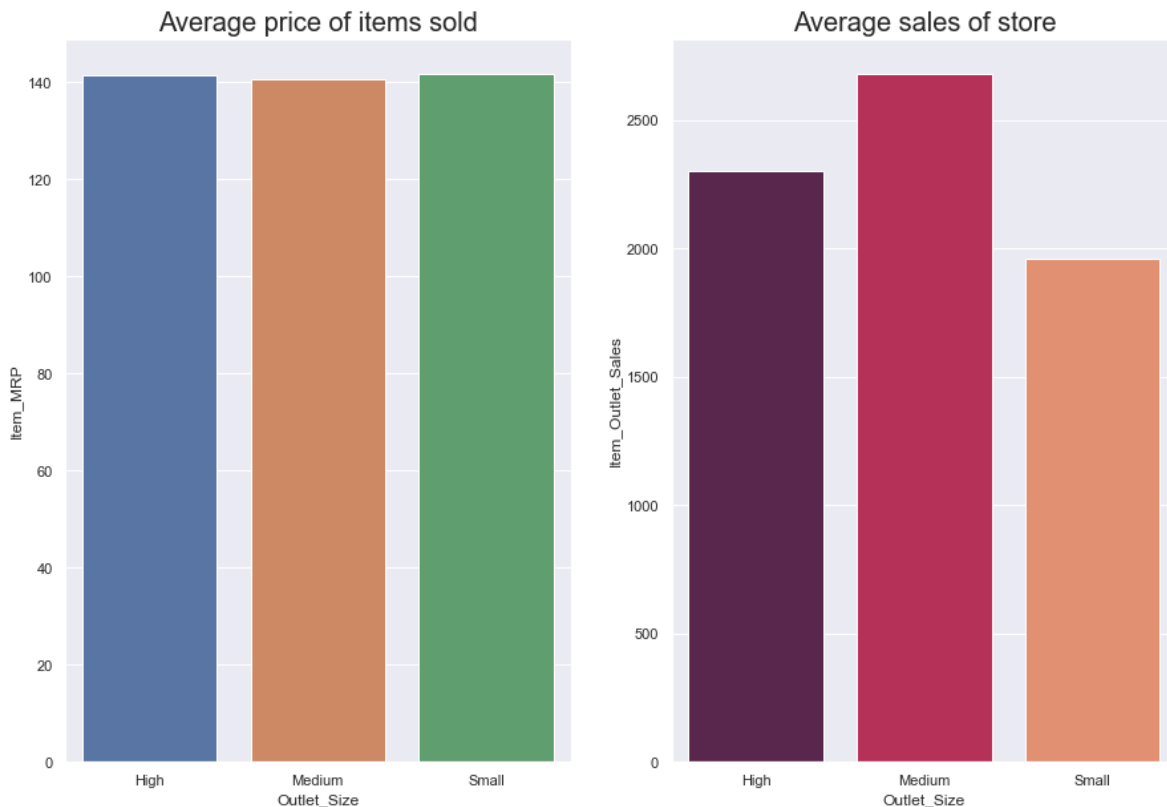
	Outlet_Size	Item_MRP	Item_Outlet_Sales
0	High	141.425982	2298.995256
1	Medium	140.590514	2681.603542
2	Small	141.756737	1960.412740

In [ ]:

```
fig8=plt.figure(figsize=(15,10))
ax1=fig8.add_subplot(121)
sns.barplot('Outlet_Size', 'Item_MRP', data=df_size_sales, ax=ax1)

ax2=fig8.add_subplot(122)
sns.barplot('Outlet_Size', 'Item_Outlet_Sales', data=df_size_sales, ax=ax2, palette='rocket')

ax1.set_title('Average price of items sold', size=20)
ax2.set_title('Average sales of store', size=20)
plt.show()
```



The average price of items sold in each outlet store size is nearly the same which is Rs 140. However, The medium stores seem to sell better followed by high sized and then small sized stores.

## 8. Outlet and Outlet location

In [ ]:

```
df_train.head()
```

Out[31]:

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier
0	FDA15	9.30	Low Fat	0.016047	Dairy	249.8092	OUT001
1	DRC01	5.92	Regular	0.019278	Soft Drinks	48.2692	OUT001
2	FDN15	17.50	Low Fat	0.016760	Meat	141.6180	OUT001
3	FDX07	19.20	Regular	0.000000	Fruits and Vegetables	182.0950	OUT001
4	NCD19	8.93	Low Fat	0.000000	Household	53.8614	OUT001

In [ ]:

```
fig9=px.sunburst(df_train,path=['Outlet_Type','Outlet_Location_Type'],color_continuous_scal  
fig9.update_layout(title='Store type with location type',title_x=0.5)  
fig9.show()
```


As we can see, majority of the stores are of type 1 supermarket distributed over various location tiers.

Supermarket type 2 and 3 are confined to only tier 3 locations. Very small section of the stores are actually grocery stores.

Let us check how do these stores sell based on location tier using a boxplot.

In [ ]:

```
plt.figure(figsize=(10,8))
sns.boxplot(y='Item_Outlet_Sales',hue='Outlet_Type',x='Outlet_Location_Type',data=df_train,
plt.show())
```



As we can see, tier 3 locations seem to be selling better than both tier 2 and tier 1. It is also to be noted that tier 3 has more number of stores in it. Hence, the sales are better too.

## Correlation heatmap

In [ ]:


```
df_train.drop('Count',axis=1,inplace=True)
```

In [ ]:

```

corrs=df_train.dropna().corr()
plt.figure(figsize=(10,8))
sns.heatmap(corrs,annot=True,fmt='.2%')
plt.show()

```



From the above, we can see that correlation of Item\_Weight is extremely low. Hence, we can simply drop this column and get done with the issues of null values. We shall similarly remove the order\_size as there is no way to deal with the null values here aswell. We would also get rid of the item\_identifier and outlet\_identifier since it is of no consequence to us.

In [ ]:

```

unn_cols=['Item_Weight','Outlet_Size','Item_Identifier','Outlet_Identifier']

for cols in unn_cols:
    df_train.drop(cols,axis=1,inplace=True)

```

## Data Wrangling



In [ ]:

```
df_train['Item_Fat_Content'].replace('Low Fat',1,inplace=True)
df_train['Item_Fat_Content'].replace('Regular',0,inplace=True)
```

In [ ]:

```
df_dummies_type=pd.get_dummies(df_train['Item_Type'])
```

In [ ]:

```
df_train=df_train.merge(df_dummies_type,on=df_train.index)
```

In [ ]:

```
df_train.drop('key_0',axis=1,inplace=True)
df_train.drop('Item_Type',axis=1,inplace=True)
```

In [ ]:

```
df_train['Outlet_Location_Type'].replace('Tier 1',1,inplace=True)
df_train['Outlet_Location_Type'].replace('Tier 2',2,inplace=True)
df_train['Outlet_Location_Type'].replace('Tier 3',3,inplace=True)
```

In [ ]:

```
df_dummies_outlet=pd.get_dummies(df_train['Outlet_Type'])
df_train=df_train.merge(df_dummies_outlet,on=df_train.index)
```

## Item Outlet Sales


In [ ]:

```
targets=df_train['Item_Outlet_Sales']
df_train.drop('Item_Outlet_Sales',axis=1,inplace=True)
df_train.head()
```

Out[43]:

	key_0	Item_Fat_Content	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Outlet_Location_Type
0	0	1	0.016047	249.8092	1999	:
1	1	0	0.019278	48.2692	2009	:
2	2	1	0.016760	141.6180	1999	:
3	3	0	0.000000	182.0950	1998	:
4	4	1	0.000000	53.8614	1987	:

5 rows × 27 columns



In [ ]: